

# dlnd\_face\_generation

May 31, 2020

## 1 Face Generation

In this project, you'll define and train a DCGAN on a dataset of faces. Your goal is to get a generator network to generate *new* images of faces that look as realistic as possible!

The project will be broken down into a series of tasks from **loading in data to defining and training adversarial networks**. At the end of the notebook, you'll be able to visualize the results of your trained Generator to see how it performs; your generated samples should look like fairly realistic faces with small amounts of noise.

### 1.0.1 Get the Data

You'll be using the [CelebFaces Attributes Dataset \(CelebA\)](#) to train your adversarial networks.

This dataset is more complex than the number datasets (like MNIST or SVHN) you've been working with, and so, you should prepare to define deeper networks and train them for a longer time to get good results. It is suggested that you utilize a GPU for training.

### 1.0.2 Pre-processed Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. Some sample data is show below.

If you are working locally, you can download this data [by clicking here](#)

This is a zip file that you'll need to extract in the home directory of this notebook for further loading and processing. After extracting the data, you should be left with a directory of data processed\_celeba\_small/

```
In [1]: # can comment out after executing
        #!unzip processed_celeba_small.zip
```

```
In [2]: data_dir = 'processed_celeba_small/'
```

```
"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

import pickle as pkl
import matplotlib.pyplot as plt
```

```
import numpy as np
import problem_unittests as tests
#import helper

%matplotlib inline
```

## 1.1 Visualize the CelebA Data

The [CelebA](#) dataset contains over 200,000 celebrity images with annotations. Since you're going to be generating faces, you won't need the annotations, you'll only need the images. Note that these are color images with [3 color channels \(RGB\)](#) each.

### 1.1.1 Pre-process and Load the Data

Since the project's main focus is on building the GANs, we've done *some* of the pre-processing for you. Each of the CelebA images has been cropped to remove parts of the image that don't include a face, then resized down to 64x64x3 NumPy images. This *pre-processed* dataset is a smaller subset of the very large CelebA data.

There are a few other steps that you'll need to **transform** this data and create a **DataLoader**.

**Exercise: Complete the following `get_dataloader` function, such that it satisfies these requirements:**

- Your images should be square, Tensor images of size `image_size x image_size` in the x and y dimension.
- Your function should return a `DataLoader` that shuffles and batches these Tensor images.

**ImageFolder** To create a dataset given a directory of images, it's recommended that you use PyTorch's [ImageFolder](#) wrapper, with a root directory `processed_celeba_small/` and data transformation passed in.

```
In [3]: # necessary imports
import torch
from torchvision import datasets
from torchvision import transforms

In [4]: def get_dataloader(batch_size, image_size, data_dir='processed_celeba_small/'):
    """
    Batch the neural network data using DataLoader
    :param batch_size: The size of each batch; the number of images in a batch
    :param img_size: The square size of the image data (x, y)
    :param data_dir: Directory where image data is located
    :return: DataLoader with batched data
    """

    # Define the transformation for size, square nd tensor transformation
    transformation = transforms.Compose([transforms.Resize(image_size),
                                         transforms.CenterCrop(image_size),
```

```

        transforms.ToTensor()])
    # ImageFolder wrapper for the folder with the content and the defined transformation
    image_data = datasets.ImageFolder(data_dir, transform=transformation)
    # Data Loader with pre-defined batch size and shuffling the batches
    data_loader = torch.utils.data.DataLoader(image_data, batch_size=batch_size, shuffle=True)
    return data_loader

```

## 1.2 Create a DataLoader

**Exercise: Create a DataLoader** `celeba_train_loader` with appropriate hyperparameters. Call the above function and create a dataloader to view images. \* You can decide on any reasonable `batch_size` parameter \* Your `image_size` **must be 32**. Resizing the data to a smaller size will make for faster training, while still creating convincing images of faces!

```

In [5]: # Define function hyperparameters
        batch_size = 128
        img_size = 32

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        # Call your function and get a dataloader
        celeba_train_loader = get_dataloader(batch_size, img_size)

```

Next, you can view some images! You should see square images of somewhat-centered faces.

Note: You'll need to convert the Tensor images into a NumPy type and transpose the dimensions to correctly display an image, suggested `imshow` code is below, but it may not be perfect.

```


In [6]: # helper display function
        def imshow(img):
            npimg = img.numpy()
            plt.imshow(np.transpose(npimg, (1, 2, 0)))

        """
        DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
        """

        # obtain one batch of training images
        dataiter = iter(celeba_train_loader)
        images, _ = dataiter.next() # _ for no labels

        # plot the images in the batch, along with the corresponding labels
        fig = plt.figure(figsize=(20, 4))
        plot_size=20
        for idx in np.arange(plot_size):
            ax = fig.add_subplot(2, plot_size/2, idx+1, xticks=[], yticks=[])
            imshow(images[idx])

```



**Exercise: Pre-process your image data and scale it to a pixel range of -1 to 1** You need to do a bit of pre-processing; you know that the output of a tanh activated generator will contain pixel values in a range from -1 to 1, and so, we need to rescale our training images to a range of -1 to 1. (Right now, they are in a range from 0-1.)

```
In [7]: # TODO: Complete the scale function
def scale(x, feature_range=(-1, 1)):
    ''' Scale takes in an image x and returns that image, scaled
        with a feature_range of pixel values from -1 to 1.
        This function assumes that the input x is already scaled from 0-1. '''
    # assume x is scaled to (0, 1)
    # scale to feature_range and return scaled x
    min, max = feature_range
    x = x * (max - min) + min
    return x
```

```
In [8]: """
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

# check scaled range
# should be close to -1 to 1
img = images[0]
scaled_img = scale(img)

print('Min: ', scaled_img.min())
print('Max: ', scaled_img.max())
```

```
Min: tensor(-0.7098)
Max: tensor(0.9294)
```

---

## 2 Define the Model

A GAN is comprised of two adversarial networks, a discriminator and a generator.

## 2.1 Discriminator

Your first task will be to define the discriminator. This is a convolutional classifier like you've built before, only without any maxpooling layers. To deal with this complex data, it's suggested you use a deep network with **normalization**. You are also allowed to create any helper functions that may be useful.

### Exercise: Complete the Discriminator class

- The inputs to the discriminator are 32x32x3 tensor images
- The output should be a single value that will indicate whether a given image is real or fake

```
In [9]: import torch.nn as nn
        import torch.nn.functional as F

In [10]: # Creating a conv helper function
        def conv(in_channels, out_channels, kernel_size=4, stride=2, padding=1, batch_norm=True):
            layers = []
            conv_layer = nn.Conv2d(in_channels, out_channels, kernel_size, stride, padding, bias=False)
            layers.append(conv_layer)

            # append conv layer
            layers.append(conv_layer)

            if batch_norm:
                # append batchnorm layer
                layers.append(nn.BatchNorm2d(out_channels))

            # using Sequential container
            return nn.Sequential(*layers)

In [11]: class Discriminator(nn.Module):

        def __init__(self, conv_dim):
            """
            Initialize the Discriminator Module
            :param conv_dim: The depth of the first convolutional layer
            """
            super(Discriminator, self).__init__()

            # complete init function
            self.conv_dim = conv_dim

            # 32x32 input
            self.conv1 = conv(3, conv_dim, batch_norm = False)
            # 16x16 input
            self.conv2 = conv(conv_dim, conv_dim*2)
            # 8x8 input
            self.conv3 = conv(conv_dim*2, conv_dim*4)
            # 4x4 out
```

```

        # fully connected layer Image: 4x4 * conv_dim*4
        self.fc = nn.Linear(conv_dim*4*4*4, 1)

    def forward(self, x):
        """
        Forward propagation of the neural network
        :param x: The input to the neural network
        :return: Discriminator logits; the output of the neural network
        """

        # Applying Relus on the 3 conv layers
        x = F.leaky_relu(self.conv1(x), 0.2)
        x = F.leaky_relu(self.conv2(x), 0.2)
        x = F.leaky_relu(self.conv3(x), 0.2)

        # Flattening
        x = x.view(-1, self.conv_dim*4*4*4)

        # FC Layer
        x = self.fc(x)

        return x

    """
    DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
    """

    tests.test_discriminator(Discriminator)

```

Tests Passed

## 2.2 Generator

The generator should upsample an input and generate a *new* image of the same size as our training data 32x32x3. This should be mostly transpose convolutional layers with normalization applied to the outputs.

### Exercise: Complete the Generator class

- The inputs to the generator are vectors of some length `z_size`
- The output should be a image of shape 32x32x3

```

In [12]: # deconv helper function
def deconv(in_channels, out_channels, kernel_size = 4, stride = 2, padding = 1, batch_norm=True):
    # create a sequence of transpose + optional batch norm layers
    layers = []
    transpose_conv_layer = nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, padding)
    # append transpose convolutional layer

```

```

layers.append(transpose_conv_layer)

if batch_norm:
    # append batchnorm layer
    layers.append(nn.BatchNorm2d(out_channels))

return nn.Sequential(*layers)

```

In [13]: `class Generator(nn.Module):`

```

def __init__(self, z_size, conv_dim):
    """
    Initialize the Generator Module
    :param z_size: The length of the input latent vector, z
    :param conv_dim: The depth of the inputs to the *last* transpose convolutional
    """
    super(Generator, self).__init__()

    self.conv_dim = conv_dim

    # Fully connected layer from z values:
    self.fc = nn.Linear(z_size, conv_dim*4*4*4)

    # Transposed convolutional layers with helper function
    self.t_conv1 = deconv(conv_dim*4, conv_dim*2)
    self.t_conv2 = deconv(conv_dim*2, conv_dim)
    self.t_conv3 = deconv(conv_dim, 3, batch_norm=False)

def forward(self, x):
    """
    Forward propagation of the neural network
    :param x: The input to the neural network
    :return: A 32x32x3 Tensor image as output
    """
    # FC Layer
    x = self.fc(x)
    # Resepahing to Tensor
    x = x.view(-1, self.conv_dim*4, 4, 4)

    # Transpose conv layers and relu activation
    x = F.relu(self.t_conv1(x))
    x = F.relu(self.t_conv2(x))

    # Last Layer and tanh activation
    x = self.t_conv3(x)
    x = F.tanh(x)

    return x

```

```

"""
DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
"""

tests.test_generator(Generator)

```

Tests Passed

## 2.3 Initialize the weights of your networks

To help your models converge, you should initialize the weights of the convolutional and linear layers in your model. From reading the [original DCGAN paper](#), they say: > All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02.

So, your next task will be to define a weight initialization function that does just this!

You can refer back to the lesson on weight initialization or even consult existing model code, such as that from [the networks.py file in CycleGAN Github repository](#) to help you complete this function.

### Exercise: Complete the weight initialization function

- This should initialize only **convolutional** and **linear** layers
- Initialize the weights to a normal distribution, centered around 0, with a standard deviation of 0.02.
- The bias terms, if they exist, may be left alone or set to 0.

```

In [14]: def weights_init_normal(m):
          """
          Applies initial weights to certain layers in a model .
          The weights are taken from a normal distribution
          with mean = 0, std dev = 0.02.
          :param m: A module or layer in a network
          """
          # classname will be something like:
          # `Conv`, `BatchNorm2d`, `Linear`, etc.
          classname = m.__class__.__name__

          if hasattr(m, 'weight') and (classname.find('Conv') != -1 or classname.find('Linear')
              m.weight.data.normal_(0, 0.02)

          # Bias term to 0
          if hasattr(m, 'bias') and m.bias is not None:
              m.bias.data.zero_()

```

## 2.4 Build complete network

Define your models' hyperparameters and instantiate the discriminator and generator from the classes defined above. Make sure you've passed in the correct input arguments.



```

In [15]: """
          DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
          """

def build_network(d_conv_dim, g_conv_dim, z_size):
    # define discriminator and generator
    D = Discriminator(d_conv_dim)
    G = Generator(z_size=z_size, conv_dim=g_conv_dim)

    # initialize model weights
    D.apply(weights_init_normal)
    G.apply(weights_init_normal)

    print(D)
    print()
    print(G)

    return D, G

```

### Exercise: Define model hyperparameters

```

In [16]: # Define model hyperparams
         d_conv_dim = 64
         g_conv_dim = 64
         z_size = 100

         """
         DON'T MODIFY ANYTHING IN THIS CELL THAT IS BELOW THIS LINE
         """

         D, G = build_network(d_conv_dim, g_conv_dim, z_size)

```

```

Discriminator(
  (conv1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
  )
  (conv2): Sequential(
    (0): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (conv3): Sequential(
    (0): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (fc): Linear(in_features=4096, out_features=1, bias=True)
)

```

```

Generator(
  (fc): Linear(in_features=100, out_features=4096, bias=True)
  (t_conv1): Sequential(

```

```

        (0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (t_conv2): Sequential(
      (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (t_conv3): Sequential(
      (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    )
  )
)

```

### 2.4.1 Training on GPU

Check if you can train on GPU. Here, we'll set this as a boolean variable `train_on_gpu`. Later, you'll be responsible for making sure that >\* Models, \* Model inputs, and \* Loss function arguments

Are moved to GPU, where appropriate.

```

In [17]: """
        DON'T MODIFY ANYTHING IN THIS CELL
        """

        import torch

        # Check for a GPU
        train_on_gpu = torch.cuda.is_available()
        if not train_on_gpu:
            print('No GPU found. Please use a GPU to train your neural network.')
        else:
            print('Training on GPU!')

```

Training on GPU!

---

## 2.5 Discriminator and Generator Losses

Now we need to calculate the losses for both types of adversarial networks.

### 2.5.1 Discriminator Losses

- For the discriminator, the total loss is the sum of the losses for real and fake images,  $d\_loss = d\_real\_loss + d\_fake\_loss$ .
- Remember that we want the discriminator to output 1 for real images and 0 for fake images, so we need to set up the losses to reflect that.

## 2.5.2 Generator Loss

The generator loss will look similar only with flipped labels. The generator's goal is to get the discriminator to *think* its generated images are *real*.

**Exercise: Complete real and fake loss functions** You may choose to use either cross entropy or a least squares error loss to complete the following `real_loss` and `fake_loss` functions.

```
In [18]: def real_loss(D_out):
    '''Calculates how close discriminator outputs are to being real.
        param, D_out: discriminator logits
        return: real loss'''
    batch_size = D_out.size(0)
    labels = torch.ones(batch_size)

    # move labels to GPU if available
    if train_on_gpu:
        labels = labels.cuda()

    # binary cross entropy with logits loss
    criterion = nn.BCEWithLogitsLoss()
    # calculate loss
    loss = criterion(D_out.squeeze(), labels)
    return loss

def fake_loss(D_out):
    '''Calculates how close discriminator outputs are to being fake.
        param, D_out: discriminator logits
        return: fake loss'''
    batch_size = D_out.size(0)
    labels = torch.zeros(batch_size)

    if train_on_gpu:
        labels = labels.cuda()

    criterion = nn.BCEWithLogitsLoss()
    # calculate loss
    loss = criterion(D_out.squeeze(), labels)
    return loss
```

## 2.6 Optimizers

**Exercise: Define optimizers for your Discriminator (D) and Generator (G)** Define optimizers for your models with appropriate hyperparameters.

```
In [19]: import torch.optim as optim

    # params
```

```

lr = 0.0002
beta1 = 0.5
beta2 = 0.999

# Create optimizers for the discriminator D and generator G
d_optimizer = optim.Adam(D.parameters(), lr, [beta1, beta2])
g_optimizer = optim.Adam(G.parameters(), lr, [beta1, beta2])

```

---

## 2.7 Training

Training will involve alternating between training the discriminator and the generator. You'll use your functions `real_loss` and `fake_loss` to help you calculate the discriminator losses.

- You should train the discriminator by alternating on real and fake images
- Then the generator, which tries to trick the discriminator and should have an opposing loss function

**Saving Samples** You've been given some code to print out some loss statistics and save some generated "fake" samples.

**Exercise: Complete the training function** Keep in mind that, if you've moved your models to GPU, you'll also have to move any model inputs to GPU.

```

In [20]: def train(D, G, n_epochs, print_every=50):
    '''Trains adversarial networks for some number of epochs
    param, D: the discriminator network
    param, G: the generator network
    param, n_epochs: number of epochs to train for
    param, print_every: when to print and record the models' losses
    return: D and G losses'''

    # move models to GPU
    if train_on_gpu:
        D.cuda()
        G.cuda()

    # keep track of loss and generated, "fake" samples
    samples = []
    losses = []

    # Get some fixed data for sampling. These are images that are held
    # constant throughout training, and allow us to inspect the model's performance
    sample_size=16
    fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
    fixed_z = torch.from_numpy(fixed_z).float()
    # move z to GPU if available

```

```

if train_on_gpu:
    fixed_z = fixed_z.cuda()

# epoch training loop
for epoch in range(n_epochs):

    # batch training loop
    for batch_i, (real_images, _) in enumerate(celeba_train_loader):

        batch_size = real_images.size(0)
        real_images = scale(real_images)

        # =====
        #          YOUR CODE HERE: TRAIN THE NETWORKS
        # =====

        # 1. Train the discriminator on real and fake images

        # Set grad to zero
        d_optimizer.zero_grad()

        # Train on GPU
        if train_on_gpu:
            real_images = real_images.cuda()

        # Apply Discriminator on the real images
        D_real = D(real_images)
        d_real_loss = real_loss(D_real)

        # Generate Fake images
        z = np.random.uniform(-1, 1, size=(batch_size, z_size))
        z = torch.from_numpy(z).float()

        # Move to GPU
        if train_on_gpu:
            z = z.cuda()

        # Apply Generator on z to generate fake images
        fake_images = G(z)

        # Running the fake images on the Discriminator
        D_fake = D(fake_images)
        d_fake_loss = fake_loss(D_fake)

        d_loss = d_real_loss + d_fake_loss
        d_loss.backward()
        d_optimizer.step()

```

```

# 2. Train the generator with an adversarial loss

# Set the Gradient to zero
g_optimizer.zero_grad()

# generate z
z = np.random.uniform(-1, 1, size=(batch_size, z_size))
z = torch.from_numpy(z).float()

# Move to GPU
if train_on_gpu:
    z = z.cuda()

# create fake images with the Generator
fake_images = G(z)

# Apply Discriminator and calculate loss
D_fake = D(fake_images)
g_loss = real_loss(D_fake)

# back-prop and optimizer
g_loss.backward()
g_optimizer.step()

# =====
#                      END OF YOUR CODE
# =====

# Print some loss stats
if batch_i % print_every == 0:
    # append discriminator loss and generator loss
    losses.append((d_loss.item(), g_loss.item()))
    # print discriminator and generator loss
    print('Epoch [{:5d}/{:5d}] | d_loss: {:.4f} | g_loss: {:.4f}'.format(
        epoch+1, n_epochs, d_loss.item(), g_loss.item()))

## AFTER EACH EPOCH##
# this code assumes your generator is named G, feel free to change the name
# generate and save sample, fake images
G.eval() # for generating samples
samples_z = G(fixed_z)
samples.append(samples_z)
G.train() # back to training mode

# Save training generator samples
with open('train_samples.pkl', 'wb') as f:

```

```

        pkl.dump(samples, f)

    # finally return losses
    return losses

```

Set your number of training epochs and train your GAN!

```

In [21]: # set number of epochs
         n_epochs = 10

```

```

"""
DON'T MODIFY ANYTHING IN THIS CELL
"""

# call training function
losses = train(D, G, n_epochs=n_epochs)

```

```

Epoch [ 1/ 10] | d_loss: 1.3848 | g_loss: 1.0317
Epoch [ 1/ 10] | d_loss: 0.0865 | g_loss: 4.0679
Epoch [ 1/ 10] | d_loss: 0.2049 | g_loss: 3.2902
Epoch [ 1/ 10] | d_loss: 0.3237 | g_loss: 3.4050
Epoch [ 1/ 10] | d_loss: 0.5329 | g_loss: 3.7378
Epoch [ 1/ 10] | d_loss: 2.0889 | g_loss: 4.2348
Epoch [ 1/ 10] | d_loss: 0.4855 | g_loss: 1.8386
Epoch [ 1/ 10] | d_loss: 0.7826 | g_loss: 1.5993
Epoch [ 1/ 10] | d_loss: 0.7872 | g_loss: 1.5260
Epoch [ 1/ 10] | d_loss: 0.6386 | g_loss: 2.4055
Epoch [ 1/ 10] | d_loss: 0.7807 | g_loss: 1.4434
Epoch [ 1/ 10] | d_loss: 0.8968 | g_loss: 0.7620
Epoch [ 1/ 10] | d_loss: 0.9718 | g_loss: 1.6253
Epoch [ 1/ 10] | d_loss: 0.9436 | g_loss: 1.3210
Epoch [ 1/ 10] | d_loss: 1.0262 | g_loss: 0.9289
Epoch [ 2/ 10] | d_loss: 0.9623 | g_loss: 1.2973
Epoch [ 2/ 10] | d_loss: 0.9433 | g_loss: 0.9757
Epoch [ 2/ 10] | d_loss: 0.9030 | g_loss: 1.1766
Epoch [ 2/ 10] | d_loss: 1.1302 | g_loss: 1.5276
Epoch [ 2/ 10] | d_loss: 1.1149 | g_loss: 1.1764
Epoch [ 2/ 10] | d_loss: 1.1557 | g_loss: 1.4719
Epoch [ 2/ 10] | d_loss: 0.8660 | g_loss: 1.3539
Epoch [ 2/ 10] | d_loss: 1.0941 | g_loss: 1.4021
Epoch [ 2/ 10] | d_loss: 1.0820 | g_loss: 1.1755
Epoch [ 2/ 10] | d_loss: 1.1008 | g_loss: 0.9431
Epoch [ 2/ 10] | d_loss: 1.0564 | g_loss: 1.4423
Epoch [ 2/ 10] | d_loss: 1.1263 | g_loss: 1.4151
Epoch [ 2/ 10] | d_loss: 1.0669 | g_loss: 1.6489
Epoch [ 2/ 10] | d_loss: 1.1321 | g_loss: 1.0313
Epoch [ 2/ 10] | d_loss: 1.0512 | g_loss: 1.5004
Epoch [ 3/ 10] | d_loss: 0.9489 | g_loss: 1.5805

```

Epoch [	3/	10]	d_loss: 0.9582	g_loss: 1.3813
Epoch [	3/	10]	d_loss: 1.2932	g_loss: 1.6055
Epoch [	3/	10]	d_loss: 1.1473	g_loss: 0.9118
Epoch [	3/	10]	d_loss: 1.1042	g_loss: 0.9264
Epoch [	3/	10]	d_loss: 1.1654	g_loss: 0.6028
Epoch [	3/	10]	d_loss: 1.0934	g_loss: 0.8811
Epoch [	3/	10]	d_loss: 0.9749	g_loss: 1.2898
Epoch [	3/	10]	d_loss: 1.1713	g_loss: 0.8371
Epoch [	3/	10]	d_loss: 0.9626	g_loss: 1.0160
Epoch [	3/	10]	d_loss: 0.8344	g_loss: 1.5120
Epoch [	3/	10]	d_loss: 1.1922	g_loss: 1.6833
Epoch [	3/	10]	d_loss: 1.3811	g_loss: 0.5569
Epoch [	3/	10]	d_loss: 1.1223	g_loss: 1.1463
Epoch [	3/	10]	d_loss: 1.0987	g_loss: 1.4966
Epoch [	4/	10]	d_loss: 1.0906	g_loss: 0.8383
Epoch [	4/	10]	d_loss: 1.0489	g_loss: 1.6813
Epoch [	4/	10]	d_loss: 0.9405	g_loss: 1.2938
Epoch [	4/	10]	d_loss: 0.8860	g_loss: 1.0859
Epoch [	4/	10]	d_loss: 1.2250	g_loss: 0.7602
Epoch [	4/	10]	d_loss: 1.0216	g_loss: 1.0612
Epoch [	4/	10]	d_loss: 0.9219	g_loss: 1.5956
Epoch [	4/	10]	d_loss: 1.0511	g_loss: 0.7719
Epoch [	4/	10]	d_loss: 0.9834	g_loss: 0.7999
Epoch [	4/	10]	d_loss: 1.4256	g_loss: 0.4232
Epoch [	4/	10]	d_loss: 1.1274	g_loss: 1.7267
Epoch [	4/	10]	d_loss: 0.9195	g_loss: 2.0237
Epoch [	4/	10]	d_loss: 1.1150	g_loss: 0.7380
Epoch [	4/	10]	d_loss: 1.1390	g_loss: 0.5704
Epoch [	4/	10]	d_loss: 0.9500	g_loss: 0.8966
Epoch [	5/	10]	d_loss: 0.9530	g_loss: 1.5228
Epoch [	5/	10]	d_loss: 1.0534	g_loss: 1.5873
Epoch [	5/	10]	d_loss: 1.3852	g_loss: 0.4986
Epoch [	5/	10]	d_loss: 0.9195	g_loss: 1.4738
Epoch [	5/	10]	d_loss: 1.0895	g_loss: 0.9799
Epoch [	5/	10]	d_loss: 0.7773	g_loss: 1.4730
Epoch [	5/	10]	d_loss: 0.9868	g_loss: 1.5019
Epoch [	5/	10]	d_loss: 0.8869	g_loss: 1.0780
Epoch [	5/	10]	d_loss: 0.9954	g_loss: 1.1470
Epoch [	5/	10]	d_loss: 0.9349	g_loss: 1.1140
Epoch [	5/	10]	d_loss: 1.0955	g_loss: 1.1504
Epoch [	5/	10]	d_loss: 0.9404	g_loss: 1.0557
Epoch [	5/	10]	d_loss: 1.0428	g_loss: 1.3134
Epoch [	5/	10]	d_loss: 0.8455	g_loss: 1.2000
Epoch [	5/	10]	d_loss: 1.1956	g_loss: 0.7354
Epoch [	6/	10]	d_loss: 0.8416	g_loss: 1.1681
Epoch [	6/	10]	d_loss: 1.0153	g_loss: 0.8833
Epoch [	6/	10]	d_loss: 1.5240	g_loss: 2.8041
Epoch [	6/	10]	d_loss: 0.7251	g_loss: 1.7810



Epoch [	6/	10]	d_loss: 1.0628	g_loss: 0.9683
Epoch [	6/	10]	d_loss: 0.9565	g_loss: 1.1104
Epoch [	6/	10]	d_loss: 1.1749	g_loss: 1.7561
Epoch [	6/	10]	d_loss: 1.0282	g_loss: 1.7698
Epoch [	6/	10]	d_loss: 1.0964	g_loss: 1.2815
Epoch [	6/	10]	d_loss: 1.1126	g_loss: 0.9641
Epoch [	6/	10]	d_loss: 1.4566	g_loss: 1.9784
Epoch [	6/	10]	d_loss: 1.0194	g_loss: 0.8733
Epoch [	6/	10]	d_loss: 0.9961	g_loss: 1.0358
Epoch [	6/	10]	d_loss: 1.6942	g_loss: 0.6681
Epoch [	6/	10]	d_loss: 0.7524	g_loss: 2.1149
Epoch [	7/	10]	d_loss: 0.9397	g_loss: 1.5187
Epoch [	7/	10]	d_loss: 1.0464	g_loss: 1.3926
Epoch [	7/	10]	d_loss: 0.8807	g_loss: 1.3250
Epoch [	7/	10]	d_loss: 0.8966	g_loss: 1.6015
Epoch [	7/	10]	d_loss: 0.9429	g_loss: 1.1847
Epoch [	7/	10]	d_loss: 0.8048	g_loss: 1.6319
Epoch [	7/	10]	d_loss: 0.9552	g_loss: 1.3203
Epoch [	7/	10]	d_loss: 1.0162	g_loss: 0.9360
Epoch [	7/	10]	d_loss: 0.9814	g_loss: 1.3738
Epoch [	7/	10]	d_loss: 0.9114	g_loss: 1.9042
Epoch [	7/	10]	d_loss: 1.0784	g_loss: 0.8970
Epoch [	7/	10]	d_loss: 0.7878	g_loss: 1.2650
Epoch [	7/	10]	d_loss: 1.0364	g_loss: 0.8628
Epoch [	7/	10]	d_loss: 0.9953	g_loss: 1.3935
Epoch [	7/	10]	d_loss: 1.2202	g_loss: 0.7879
Epoch [	8/	10]	d_loss: 1.1796	g_loss: 1.2244
Epoch [	8/	10]	d_loss: 1.1412	g_loss: 0.9225
Epoch [	8/	10]	d_loss: 0.9066	g_loss: 1.3896
Epoch [	8/	10]	d_loss: 1.0572	g_loss: 0.7858
Epoch [	8/	10]	d_loss: 1.0933	g_loss: 0.9334
Epoch [	8/	10]	d_loss: 1.3175	g_loss: 0.9971
Epoch [	8/	10]	d_loss: 1.0659	g_loss: 1.0750
Epoch [	8/	10]	d_loss: 0.9292	g_loss: 1.6820
Epoch [	8/	10]	d_loss: 1.3219	g_loss: 1.0629
Epoch [	8/	10]	d_loss: 0.9993	g_loss: 1.3302
Epoch [	8/	10]	d_loss: 0.9065	g_loss: 1.3843
Epoch [	8/	10]	d_loss: 0.9728	g_loss: 1.1199
Epoch [	8/	10]	d_loss: 0.9195	g_loss: 1.7747
Epoch [	8/	10]	d_loss: 0.7410	g_loss: 1.6200
Epoch [	8/	10]	d_loss: 0.9098	g_loss: 0.7428
Epoch [	9/	10]	d_loss: 1.7157	g_loss: 3.1344
Epoch [	9/	10]	d_loss: 0.9292	g_loss: 1.3809
Epoch [	9/	10]	d_loss: 0.9104	g_loss: 0.9060
Epoch [	9/	10]	d_loss: 0.8020	g_loss: 1.3174
Epoch [	9/	10]	d_loss: 0.9076	g_loss: 1.0880
Epoch [	9/	10]	d_loss: 0.9653	g_loss: 0.7784
Epoch [	9/	10]	d_loss: 0.7974	g_loss: 1.1021

```

Epoch [ 9/ 10] | d_loss: 0.9970 | g_loss: 0.9002
Epoch [ 9/ 10] | d_loss: 0.9096 | g_loss: 1.2509
Epoch [ 9/ 10] | d_loss: 1.0104 | g_loss: 1.6557
Epoch [ 9/ 10] | d_loss: 1.0152 | g_loss: 0.9667
Epoch [ 9/ 10] | d_loss: 1.0969 | g_loss: 1.6327
Epoch [ 9/ 10] | d_loss: 0.9415 | g_loss: 1.3703
Epoch [ 9/ 10] | d_loss: 0.8479 | g_loss: 2.1367
Epoch [ 9/ 10] | d_loss: 1.1096 | g_loss: 0.8890
Epoch [ 10/ 10] | d_loss: 0.8526 | g_loss: 1.3131
Epoch [ 10/ 10] | d_loss: 1.0300 | g_loss: 0.6714
Epoch [ 10/ 10] | d_loss: 0.8223 | g_loss: 1.0999
Epoch [ 10/ 10] | d_loss: 0.9820 | g_loss: 1.7601
Epoch [ 10/ 10] | d_loss: 0.8686 | g_loss: 1.7095
Epoch [ 10/ 10] | d_loss: 0.9627 | g_loss: 0.8961
Epoch [ 10/ 10] | d_loss: 1.0138 | g_loss: 1.6701
Epoch [ 10/ 10] | d_loss: 0.8379 | g_loss: 1.3998
Epoch [ 10/ 10] | d_loss: 1.0188 | g_loss: 1.0016
Epoch [ 10/ 10] | d_loss: 1.0127 | g_loss: 0.8766
Epoch [ 10/ 10] | d_loss: 0.9614 | g_loss: 1.4913
Epoch [ 10/ 10] | d_loss: 0.8748 | g_loss: 1.2173
Epoch [ 10/ 10] | d_loss: 0.9916 | g_loss: 1.5193
Epoch [ 10/ 10] | d_loss: 0.8429 | g_loss: 1.6373
Epoch [ 10/ 10] | d_loss: 0.9534 | g_loss: 0.9294

```

## 2.8 Training loss

Plot the training losses for the generator and discriminator, recorded after each epoch.

```


In [22]: fig, ax = plt.subplots()
         losses = np.array(losses)
         plt.plot(losses.T[0], label='Discriminator', alpha=0.5)
         plt.plot(losses.T[1], label='Generator', alpha=0.5)
         plt.title("Training Losses")
         plt.legend()

```

```

Out[22]: <matplotlib.legend.Legend at 0x7f531c3e6978>

```




## 2.9 Generator samples from training

View samples of images from the generator, and answer a question about the strengths and weaknesses of your trained models.

```
In [23]: # helper function for viewing a list of passed in sample images
def view_samples(epoch, samples):
    fig, axes = plt.subplots(figsize=(16,4), nrows=2, ncols=8, sharey=True, sharex=True)
    for ax, img in zip(axes.flatten(), samples[epoch]):
        img = img.detach().cpu().numpy()
        img = np.transpose(img, (1, 2, 0))
        img = ((img + 1)*255 / (2)).astype(np.uint8)
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        im = ax.imshow(img.reshape((32,32,3)))

In [24]: # Load samples from generator, taken while training
with open('train_samples.pkl', 'rb') as f:
    samples = pkl.load(f)

In [25]: _ = view_samples(-1, samples)
```



### 2.9.1 Question: What do you notice about your generated samples and how might you improve this model?

When you answer this question, consider the following factors: \* The dataset is biased; it is made of "celebrity" faces that are mostly white \* Model size; larger models have the opportunity to learn more features in a data feature space \* Optimization strategy; optimizers and number of epochs affect your final result

#### Answer:

I am quite surprised by the output of the network. Most of the images look like true faces. I do see some slight glitches, even though those could still be real people.

All my faces in that sample seem to be good in detecting the features of a normal face, hence my convolutional network is doing a fine job in feature detection.

It seems to only clash when it is about face vs. background and face vs. hair/hat. That seems to be the areas for error.

Obviously, we could deepen the network by adding another layer or adjusting the depth of the layers. Since the feature detection is good, I would be inclined to rather work on changing the training to more epochs, increase the batch\_size to 256, I also could try out a smaller learning rate as well as smoothing.

### 2.9.2 Submitting This Project

When submitting this project, make sure to run all the cells before saving the notebook. Save the notebook file as "dlnd\_face\_generation.ipynb" and save it as a HTML file under "File" -> "Download as". Include the "problem\_unittests.py" files in your submission.